

Statistical Models for Repeated Categorical Ratings: The R Package *rater*

by Jeffrey M. Pullin, Lyle C. Gurrin, and Damjan Vukcevic

Abstract A common problem in many disciplines is the need to assign a set of items into categories or classes with known labels. This is often done by one or more expert raters, or sometimes by an automated process. If these assignments or ‘ratings’ are difficult to make accurately, a common tactic is to repeat them by different raters, or even by the same rater multiple times on different occasions. We present an R package *rater*, available on CRAN, that implements Bayesian versions of several statistical models for analysis of repeated categorical rating data. Inference is possible for the true underlying (latent) class of each item, as well as the accuracy of each rater. The models are extensions of, and include, the Dawid–Skene model, and we implemented them using the Stan probabilistic programming language. We illustrate the use of *rater* through a few examples. We also discuss in detail the techniques of marginalisation and conditioning, which are necessary for these models but also apply more generally to other models implemented in Stan.

1 Introduction

The practice of measuring phenomena by having one or more raters assign to items one of a set of ratings is common across many fields. For example, in medicine one or more doctors may classify a diagnostic image as being evidence for one of several diagnoses, such as types of cancer. This process of rating is often done by one or more expert raters, but may also be performed by a large group of non-experts (e.g., in natural language processing (NLP), where a large number of crowd-sourced raters are used to classify pieces of text; see Passonneau and Carpenter (2014) and Ipeirotis, Provost, and Wang (2010)) or by a machine or other automated process (e.g., a laboratory diagnostic test, where ‘test’ may refer to either the raters and/or the ratings). Other fields where ratings occur include astronomy, for classifying astronomical images (Smyth et al. 1994), and bioinformatics, for inferring error rates for some bioinformatics algorithms (Jakobsdottir and Weeks 2007). Indeed, both studies apply the Dawid–Skene model that we present below and implement in our software package.

The assignment of items to categories, which are variously referred to as gradings, annotations or labelled data, we will call *ratings*. Our hope is that these are an accurate reflection of the true characteristics of the items being rated. However, this is not guaranteed; all raters and rating systems are fallible. We would expect some disagreement between the raters, especially from non-experts, and even expert raters may disagree when there are some items that are more difficult to rate than others. A typical way of dealing with this problem is to obtain multiple ratings for each item from different raters, some of whom might rate the item more than once on separate occasions. By using an aggregate index (e.g., by averaging the ratings) one can hope to reduce bias and variance when estimating both population frequencies and individual item categories.

Despite the fact that ratings data of this type are common, there are few software packages that can be used to fit useful statistical models to them. To address this, we introduce *rater*, a software package for R (R Core Team 2021) designed specifically for such data and available on CRAN. Our package provides the ability to fit the Dawid–Skene model (Dawid and Skene 1979) and several of its extensions. The goal is to estimate accurately the underlying true class of each item (to the extent that this is meaningful and well-defined) as well as to quantify and characterise the accuracy of each rater. Models of these data should account for the possibility that accuracy differs between the raters and that the errors they make might be class-specific. The package accepts data in multiple formats and provides functions for extracting and visualizing parameter estimates.

While our package implements optimisation-based inference, the default mode of inference is Bayesian, using Markov Chain Monte Carlo (MCMC). We chose a Bayesian approach for several reasons. The first is because the standard classical approach, using an EM algorithm, often gives parameter estimates on the boundary of the parameter space (i.e., probabilities of exactly 0 or 1), which would typically be implausible. Examples of this can be seen in the original paper by Dawid and Skene (1979), and also in our own results below that compare MCMC and optimisation-based inference (e.g., compare Figure 2 and Figure 4). A second reason is to take advantage of the convenient and powerful Stan software (Stan Development Team 2021), which made implementing our package considerably easier. Stan is a probabilistic programming language that implements MCMC and optimisation algorithms for fitting Bayesian models. The primary inference algorithm used by Stan is the No U-Turn Sampler (NUTS) variant of the Hamiltonian Monte Carlo (HMC) sampler, which uses gradient information to guide exploration of the posterior distribution. A third reason is that

Table 1: Formats for balanced rating data.

(a) Long			(b) Wide			(c) Wide (grouped)		
Item	Rater	Rating	Item	Rater 1	Rater 2	Rater 1	Rater 2	Tally
1	1	3	1	3	4	3	4	1
1	2	4	2	2	2	2	2	2
2	1	2	3	2	2			
2	2	2						
3	1	2						
3	2	2						

Table 2: Formats for unbalanced rating data. Missing values are indicated by 'NA'.

(a) Long			(b) Wide			(c) Wide (grouped)		
Item	Rater	Rating	Item	Rater 1	Rater 2	Rater 1	Rater 2	Tally
1	1	3	1	3	4	3	4	1
1	2	4	2	2	2	2	2	2
2	1	2	3	2	2	2	2	2
2	2	2	4	3	NA	3	NA	2
3	1	2	5	3	NA	NA	4	1
3	2	2	6	NA	4			
4	1	3						
5	1	3						
6	2	4						

it allows users to easily incorporate any prior information when it is available. In this context, prior information might be information about the quality of the raters or specific ratings mistakes that are more common.

We illustrate our methods and software using two examples where the raters are healthcare practitioners: one with anaesthetists rating the health of patients, and another with dentists assessing the health of teeth. In both examples, we are interested in obtaining more accurate health assessments by combining all of the experts' ratings rather than using the ratings of a single expert, as well as estimating the accuracy of each expert and characterising the types of errors they make.

Our paper is organised as follows. [Section 2](#) describes the various data formats that can be used with the function implemented in the package. [Section 4](#) introduces the Dawid-Skene model and several of its extensions. [Section 5](#) briefly describes the implementation and goes into more detail about some of the more interesting aspects, followed by an exposition of the technique called Rao-Blackwellization in [Section 6](#), which is connected with many aspects of the implementation. [Section 7](#) introduces the user interface of [rater](#) along with some examples. [Section 8](#) concludes with a discussion of the potential uses of the model.

2 Data formats

There are several different formats in which categorical rating data can be expressed. We make a distinction between *balanced* and *unbalanced* data. Balanced data have each item rated at least once by each rater. Unbalanced data contain information on some items that were not rated by every rater; stated alternatively, there is at least one rater who does not rate every item.

[Table 1](#) and [Table 2](#) illustrate three different formats for categorical rating data. Format (a) has one row per item-rater rating episode. It requires variables (data fields) to identify the item, the rater and the rating. This presentation of data is often referred to as the 'long' format. It is possible for a rater to rate a given item multiple times; each such rating would have its own row in this format. Format (b) has one row per item, and one column per rater, with each cell being a single rating. Format (c) aggregates rows (items) that have identical rating patterns and records a frequency tally. Formats (b) and (c) are examples of 'wide' formats. They only make sense for data where each rater rates each item at most once. Note that the long format will never require any structural missing values for any of the variables, whereas the wide formats require the use of a missing value indicator unless they represent data generated by a balanced design, see [Table 2](#). One benefit of the grouped format is that it allows a more computationally efficient implementation of the likelihood function. See [Section 5](#) for details about how this is implemented in [rater](#).

3 Existing approaches

A typical approach for modelling categorical ratings is to posit an unobserved underlying true class for each item, i.e., we represent the true class as a latent variable. One of the first such models was developed by Dawid and Skene (1979). This model is still studied and used to this day, and has served as a foundation for several extensions (e.g., Paun et al. 2018). Our R package, **rater**, is the first one specifically designed to provide inference for the Dawid–Skene model and its variants.

The Python library **pyanno** (Berkes et al. 2011) also implements a Bayesian version of the Dawid–Skene model in addition to the models described by Rzhetsky (2009). However, unlike **rater**, **pyanno** only supports parameter estimation via optimisation rather than MCMC; i.e., it will only compute posterior modes rather than provide samples from the full posterior distribution. In addition, **pyanno** does not support variants of the Dawid–Skene model nor the support for grouped data implemented in **rater**.

More broadly, many different so-called ‘latent class’ models have been developed and implemented in software, for a wide diversity of applications (e.g., Goodman 1974). A key aspect of the categorical ratings context is that we assume our categories are known ahead of time, i.e., the possible values for the latent classes are fixed. The Dawid–Skene model has this constraint built-in, whereas other latent class models are better tailored to other scenarios. Nevertheless, many of these models are closely related and are implemented as R packages.

Of most interest for categorical ratings analysis are the packages **poLCA** (Linzer and Lewis 2011), **BayesLCA** (White and Murphy 2014) and **randomLCA** (Beath 2017), all of which are capable of fitting limited versions of the Dawid–Skene model. We explore the relationship between different latent class models in more detail in Section 4.7. Briefly: **randomLCA** can fit the Dawid–Skene model only when the data uses binary categories, **BayesLCA** can only fit the homogeneous Dawid–Skene model, and **poLCA** can fit the Dawid–Skene model with an arbitrary number of categories but only supports wide data (where raters do not make repeated ratings on items). Neither **poLCA** nor **randomLCA** support fully Bayesian inference with MCMC, and none of the packages support fitting variants of the Dawid–Skene model (which are available in **rater**).

4 Models

One of the first statistical models proposed for categorical rating data was that of Dawid and Skene (1979). We describe this model below, extended to include prior distributions to allow for Bayesian inference. Recently, a number of direct extensions to the Dawid–Skene model have been proposed. We describe several of these below, most of which are implemented in **rater**. For ease of exposition, unless otherwise stated, all notation in this section will assume we are working with balanced data and where each item is rated exactly once by each rater, one exception is that we present the Dawid–Skene model for arbitrary data at the end of Section 4.1.

Assume we have I items (for example, images, people, etc.) and J raters, with each item presumed to belong to one of the K categories: we refer to this as its ‘true’ category and also as its latent class (since it is unobserved). Let $y_{i,j} \in \{1, \dots, K\}$ be the rating for item i given by rater j .

4.1 Dawid–Skene model

The model has two sets of parameters:

- π_k : the prevalence of category k in the population from which the items are sampled. $\pi_k \in (0, 1)$ for $k \in \{1, \dots, K\}$, with $\sum_{k=1}^K \pi_k = 1$. All classes have some non-zero probability of occurring. We collect all the π_k parameters together into a vector, $\pi = (\pi_1, \dots, \pi_K)$.
- $\theta_{j,k,k'}$: the probability that rater j responds with class k' when rating an item of true class k . Here, $j \in \{1, \dots, J\}$ and $k, k' \in \{1, 2, \dots, K\}$. We will refer to the $K \times K$ matrix $\theta_{j,k,k'}$ for a given j as the *error matrix* for rater j . We represent the k th row of this matrix as the vector $\theta_{j,k} = (\theta_{j,k,1}, \theta_{j,k,2}, \dots, \theta_{j,k,K})$: this shows how rater j responds to an item with true class k , by rating it as being in class $1, 2, \dots, K$, according to the respective probabilities.

The model also has the following set of latent variables:

- z_i : the true class of item i . $z_i \in \{1, \dots, K\}$ for $i \in \{1, \dots, I\}$.

Under the Bayesian perspective, the model parameters and latent variables are both unobserved random variables that define the joint probability distribution of the model. As such, inference on

them is conducted in the same way. However, we make the distinction here to help clarify aspects of our implementation later on.

The model is defined by the following distributional assumptions:

$$\begin{aligned} z_i &\sim \text{Categorical}(\pi), \quad \forall i \in \{1, \dots, I\}, \\ y_{i,j} \mid z_i &\sim \text{Categorical}(\theta_{j,z_i}), \quad \forall i \in \{1, \dots, I\}, j \in \{1, \dots, J\}. \end{aligned}$$

In words, the rating for item i given by rater j will follow the distribution specified by: the error matrix for rater j , but taking only the row of that matrix that corresponds to the **value** of the latent variable for item i (the z_i th row).

We now have a fully specified model, which allows use of likelihood-based methods. The likelihood function for the observed data (the $y_{i,j}$ s) is

$$\Pr(y \mid \theta, \pi) = \prod_{i=1}^I \left(\sum_{k=1}^K \left(\pi_k \cdot \prod_{j=1}^J \theta_{j,k,y_{i,j}} \right) \right).$$

Note that the unobserved latent variables, $z = (z_1, z_2, \dots, z_I)$, do not appear because they are integrated out (via the sum over the categories k). Often it is useful to work with the complete data likelihood where the z have specific values, for example if implementing an EM algorithm (such as described by Dawid and Skene (1979)). The somewhat simpler likelihood function in that case is

$$\Pr(y, z \mid \theta, \pi) = \prod_{i=1}^I \left(\pi_{z_i} \cdot \prod_{j=1}^J \theta_{j,z_i,y_{i,j}} \right).$$

In our implementation we use the first version of the likelihood, which marginalises over z (see [Section 5.1](#)). We do this to avoid needing to sample from the posterior distribution of z using Stan, as the HMC sampling algorithm used by Stan requires gradients with respect to all parameters and gradients cannot be calculated for discrete parameters such as z . Alternative Bayesian implementations, perhaps using Gibbs sampling approaches, could work with the second version of the likelihood and sample values of z directly.

To allow for Bayesian inference, we place weakly informative prior probability distributions on the parameters:

$$\begin{aligned} \pi &\sim \text{Dirichlet}(\alpha), \\ \theta_{j,k} &\sim \text{Dirichlet}(\beta_k). \end{aligned}$$

The hyper-parameters defining these prior probability distributions are:

- α : a vector of length K with all elements greater than 0
- β : a $K \times K$ matrix with all elements greater than 0, with β_k referring to the k th row of this matrix.

We use the following default values for these hyper-parameters in **rater**: $\alpha_k = 3$ for $k \in \{1, \dots, K\}$, and

$$\beta_{k,k'} = \begin{cases} Np & \text{if } k = k' \\ \frac{N(1-p)}{K-1} & \text{otherwise} \end{cases} \quad \forall k, k' \in \{1, \dots, K\}.$$

Here, N corresponds to an approximate pseudocount of hypothetical prior observations and p an approximate probability of a correct rating (applied uniformly across all raters and categories). This affords us the flexibility to centre the prior on a particular assumed accuracy (via p) as well as tune how informative the prior will be (via N). In **rater** we set the default values to be $N = 8$ and $p = 0.6$, reflecting a weakly held belief that the raters are have a better than coin-flip chance of choosing the correct class. This should be suitable for many datasets, where the number of categories is small, for example, less than ten. This default would, however, be optimistic if the number of categories is very large, for example, one hundred. In that case it would make sense for the user to specify a more realistic prior (see [Section 7.5](#)). A derivation of the default prior specification is shown in [Section 4.2](#).

We can also write the Dawid–Skene model in notation that does not assume the data are balanced. Let I, J, K be as above. Let N be the total number of ratings in the dataset (previously we had the special case $N = I \cdot J$). Define the following quantities relating to the n th rating: it was performed by rater j_n , who rated item i_n , and the rating itself is y_n . We can now write the Dawid–Skene model as:

$$\begin{aligned} z_i &\sim \text{Categorical}(\pi), \quad \forall i \in \{1, \dots, I\}, \\ y_n \mid z_{i_n} &\sim \text{Categorical}(\theta_{j_n, z_{i_n}}), \quad \forall n \in \{1, \dots, N\}. \end{aligned}$$

4.2 Hyper-parameters for the Dawid–Skene model

Hyper-parameters for the error matrices

The values we proposed for the β hyper-parameters are designed to be flexible and have an approximate intuitive interpretation. The derivation is as follows.

First, consider the variant of the model where the true latent class (z) of each item is known; this model is described under ‘Case 1. True responses are available’ by Dawid and Skene (1979). Under this model, we can ignore π because the distribution of y depends only on θ . It is a categorical distribution with parameters given by specific rows of the error matrices (as determined by z , which is known) rather than being a sum over all possible rows (according to the latent class distribution).

Under this model, the Dirichlet prior is conjugate; we obtain a Dirichlet posterior for each row of each rater’s error matrix. Consider an arbitrary rater j and an arbitrary latent class k' . Let $c = (c_1, c_2, \dots, c_K)$ be the number of times this rater rates an item as being in each of the K categories when it is of true class k' . Let $n_{jk'} = \sum_{k=1}^K c_k$ be the total number of such ratings. Also, recall that $\theta_{j,k'}$ refers to the vector of probabilities from the k' th row of the error matrix for rater j . We have that

$$c \sim \text{Multinomial}(n_{jk'}, \theta_{j,k'}),$$

giving the posterior

$$\theta_{j,k'} \mid c \sim \text{Dirichlet}(\beta_{k'} + c).$$

Under this model, the hyper-parameter vector $\beta_{k'}$ has the same influence on the posterior as the vector of counts c . We can therefore interpret the hyper-parameters as ‘pseudocounts’. This gives a convenient way to understand and set the values of the hyper-parameters. Our choices were as follows:

- We constrain their sum to be a given value $N = \sum_{k=1}^K \beta_{k',k}$, which we interpret as a pseudo-sample size (of hypothetical ratings from rater j of items with true class k').
- We then set $\beta_{k',k'}$ to a desired value so that it reflects a specific assumed accuracy. In particular, the prior mean of $\theta_{k',k'}$, the probability of a correct rating, will be $\beta_{k',k'} / N$. Let this be equal to p , an assumed average prior accuracy.
- Finally, in the absence of any information about the different categories, it is reasonable to treat them as exchangeable and thus assume that errors in the ratings are on average uniform across categories. This implies that the values of all of the other hyper-parameters ($\beta_{k',k}$ where $k \neq k'$) are equal and need to be set to a value to meet the constraint that the sum of the hyper-parameters in the vector $\beta_{k'}$ is N .
- These choices imply the hyper-parameter values described in [Section 4.1](#).

In the Dawid–Skene model, the latent classes (z) are of course not known. Therefore, we do not have a direct interpretation of the β hyper-parameters as pseudocounts. However, we can treat them approximately as such.

We chose $N = 8$ and $p = 0.6$ as default values based on obtaining reasonably good performance in simulations across a diverse range of scenarios (data not shown).

Relationship to previous work

The *Stan User’s Guide* (Section 7.4) (Stan Development Team 2021) suggests the following choice of values for β :

$$\beta_{k',k} = \begin{cases} 2.5 \times K & \text{if } k' = k \\ 1 & \text{otherwise} \end{cases}, \quad \forall k', k \in \{1, \dots, K\}.$$

It is interesting to compare this to our suggested values, above. By equating the two, we can re-write the Stan choice in terms of p and N :

$$\begin{aligned} p &= 2.5 \times K / N \\ N &= 3.5 \times K - 1. \end{aligned}$$

We see that the Stan default is to use an increasingly informative prior (N gets larger) as the number of categories (K) increases. The assumed average prior accuracy also varies based on K . For example:

- For $K = 2$ categories, $p = 0.83$ and $N = 6$.
- For $K = 5$ categories, $p = 0.76$ and $N = 16.5$.
- As K increases, p slowly reduces, with a limiting value of $2.5/3.5 = 0.71$.

Our default values of $p = 0.6$ and $N = 8$ give a prior that is less informative and less optimistic about the accuracy of the raters.

In practice we have experienced issues when fitting models via optimisation when the non-diagonal entries of β are less than 1 (i.e., when $\beta_{k',k} < 1$ for some $k' \neq k$). The default hyper-parameters selected in [rater](#) ensure that when there are only a few possible categories—specifically, when $K \leq 4$ —then this does not occur. When K is larger than this and the default hyper-parameters give rise to non-diagonal entries smaller than 1, [rater](#) will report a useful warning message. This only arises when using optimisation; it does not arise when using MCMC. Users who wish to use optimisation with a large number of categories can specify different hyper-parameters to avoid the warning message.

Hyper-parameters for the prevalence

The α hyper-parameters define the prior for the prevalence parameters π . All of these are essentially nuisance parameters and not of direct interest, and would typically have little influence on the inference for θ or z . The inclusion in the model of a large amount of prior information on the values of the prevalence parameters would clearly influence the posterior distributions of the corresponding population class frequencies. We would, however, expect inferences about the accuracy of raters (whether this is assumed to be class-specific or not) to depend largely on the number of ratings captured by the dataset, not the prevalence parameters themselves. We have, therefore, not explored varying α and have simply set them to the same values as suggested in the *Stan User's Guide* (Stan Development Team 2021): $\alpha_k = 3$ for all k .

4.3 Hierarchical Dawid–Skene model

Paun et al. (2018) introduce a ‘Hierarchical Dawid–Skene’ model that extends the original one by replacing the Dirichlet prior on $\theta_{j,k}$ (and the associated hyper-parameter β) with a hierarchical specification that allows for partial pooling, which allows the sharing of information about raters’ performance across rater-specific parameters. It requires two $K \times K$ matrices of parameters, μ and σ , with

$$\begin{aligned} \mu_{k,k'} &\sim \begin{cases} \text{Normal}(2, 1), & k = k' \\ \text{Normal}(0, 1), & k \neq k' \end{cases} \quad \forall k, k' \in \{1, \dots, K\} \\ \sigma_{k,k'} &\sim \text{Half-Normal}(0, 1) \quad \forall k, k' \in \{1, \dots, K\}. \end{aligned}$$

We then have that:

$$\gamma_{j,k,k'} \sim \text{Normal}(\mu_{k,k'}, \sigma_{k,k'}) \quad \forall j \in \{1, \dots, J\}, k, k' \in \{1, \dots, K\}$$

which are then normalised into probabilities via the multinomial logit link function (also known as the *softmax* function) in order to define the elements of the error matrices, which are:

$$\theta_{j,k,k'} = \frac{e^{\gamma_{j,k,k'}}}{\sum_{k'=1}^K e^{\gamma_{j,k,k'}}}.$$

Other details, such as the distribution of z and y , are defined in the same way as per the Dawid–Skene model. The implementation in [rater](#) differs from the implementation from Paun et al. (2018) by modifying the prior parameters to encode the assumption that the raters are generally accurate. Specifically, the higher means of the diagonal elements of μ encode that the raters are accurate, as after transformation by the softmax function larger values of μ will produce higher probabilities in θ . The assumption that the raters are accurate is also encoded in the prior distribution for θ in the Dawid–Skene model described above.

One interpretation of the hierarchical model is as a ‘partial pooling’ compromise between the full Dawid–Skene model and the model with only one rater (see [Section 4.5](#)). This is depicted in [Figure 1](#). Another interpretation of the model is that it treats the accuracy of raters given a specific latent class as a random effect, not a fixed effect as in the Dawid–Skene model. Paun et al. (2018) show, via a series of simulations and examples, that this hierarchical model generally improves the quality of predictions over the original Dawid–Skene model.

4.4 Class-conditional Dawid–Skene

Another variant of the Dawid–Skene model imposes constraints on the entries in the error matrices of the raters. Given the presumption of accurate raters, one natural way to do this is to force all non-diagonal entries in a given row of the error matrix to take a constant value that is smaller than the

corresponding diagonal entry. Formally the model is the same as the Dawid–Skene model except that we have:

$$\theta_{j,k,k'} = \begin{cases} p_{j,k} & \text{if } k = k' \\ \frac{1 - p_{j,k}}{(K - 1)} & \text{otherwise.} \end{cases}$$

To make the model fully Bayesian we use the following prior probability distributions:

$$p_{j,k} \sim \text{Beta}(\beta_{1,k}, \beta_{2,k}), \quad \forall j \in 1, \dots, J.$$

In **rater** we set $\beta_{1,k} = Np$ and $\beta_{2,k} = N(1 - p)$ for all k , where $N = 8$ and $p = 0.6$ as in [Section 4.1](#). While the constraints in this model may lead to imprecise or biased estimation of parameters for any raters that are substantially inaccurate or even antagonistic, the reduced number of parameters of this model relative to the full Dawid–Skene model make it much easier to fit. This model was referred to as the ‘class-conditional Dawid–Skene model’ by Li and Yu (2014).

4.5 Homogeneous Dawid–Skene model

Another related model is the ‘multinomial model’ described by Paun et al. (2018), where the full Dawid–Skene model is constrained so that the accuracy of all raters is assumed to be identical. The constraint can be equivalently formulated as the assumption that all of the ratings were done by a single rater, thus the raters and the sets of rating they produce are exchangeable. This model can be fitted using **rater** by simply modifying the data so that only one rater rates every item present and then fitting the usual Dawid–Skene model.

4.6 Relationships between the models

All of the models presented in this paper can be directly related to the original Dawid–Skene model. [Figure 1](#) shows the relationships between the models implemented in the package, which are coloured blue. The hierarchical model can be seen as a ‘partial pooling’ model where information about the performance of the raters is shared. The Dawid–Skene model can then be seen as the ‘no pooling’ extreme where all raters’ error matrices are estimated independently, while the homogeneous model is the other extreme of complete pooling where one error matrix is fitted for all of the raters. In addition, the class-conditional model is a direct restriction of the Dawid–Skene model where extra constraints are placed on the error matrices.

4.7 Relationships to existing models

The Dawid–Skene model is also closely related to existing latent class models. Usually, in latent class models, the number of latent classes can be chosen to maximise a suitable model selection criterion, such as the BIC, or to make the interpretation of the latent classes easier. In the Dawid–Skene model, however, the number of latent classes must be fixed to the number of categories to ensure that the error matrices can be interpreted in terms of the ability of the raters to rate the categories that appear in the data. Some specific models to which the Dawid–Skene model is related are depicted in [Figure 1](#) and are coloured white.

The latent class model implemented in **BayesLCA** is equivalent to the Dawid–Skene model fitted to binary data if the number of classes (G in the notation of White and Murphy (2014)) is set to 2. In that case, each of the M dimensions of the binary response variable can be interpreted as a rater, so that $J = M$.

The Dawid–Skene model is also a special case of the model implemented in **polCA** (Linzer and Lewis 2011). The models are equivalent if they have the same number of latent classes ($R = K$, where R is the number of latent classes in the **polCA** model) and if all of the J categorical variables in the **polCA** model take values in the set $\{1, \dots, K\}$ (that is, $K_j = K$ for all $j \in J$), so that each of the J variables can be interpreted as the ratings of a particular rater.

In addition, the Dawid–Skene model is a special case of the model implemented in **randomLCA** (Beath 2017). When no random effects are specified and $K = 2$, the model is equivalent to the Dawid–Skene model with two classes. While **randomLCA** allows selecting an arbitrary number of latent classes, it only supports binary data.

Finally, when $K = 2$ the Dawid–Skene model reduces to the so-called ‘traditional’ or ‘standard’ latent class model as first described by Hui and Walter (1980). See Asselineau et al. (2018) for a more modern description.

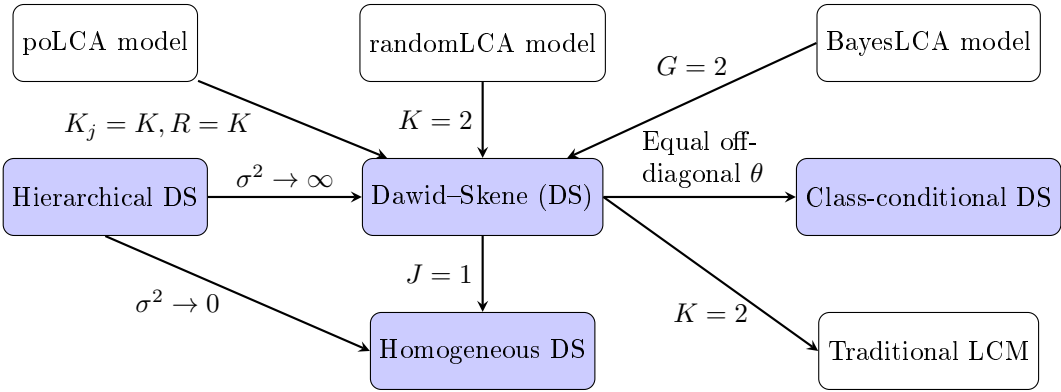


Figure 1: Relationships between models. Models coloured blue are implemented in [rater](https://CRAN.R-project.org/package=rater). DS: Dawid-Skene. LCM: latent class model.

Table 3: Features of [rater](#) and existing R packages for fitting latent class models. DS: Dawid-Skene.

Package	DS model	Fitting method	Response type	Repeated ratings	Data formats
rater	Yes	Bayesian	Polytomous	Yes	Wide, grouped, long
BayesLCA	When $K = 2$	Bayesian	Binary	No	Wide, grouped
randomLCA	When $K = 2$	Frequentist	Binary	No	Wide, grouped
poLCA	Yes	Frequentist	Polytomous	No	Wide

4.8 Relationships to existing packages

These relationships to existing latent class models means the functionality implemented in [rater](#) is similar to that of existing R packages for fitting these models. The features of [rater](#) and existing packages are summarised in [Table 3](#).

Some further details about these relationships:

- The two Bayesian packages implement different methods of inference. [rater](#) uses Hamiltonian Monte Carlo and optimisation of the log-posterior. [BayesLCA](#) implements Gibbs sampling, the expectation–maximization (EM) algorithm and a Variational Bayesian approach.
- Unlike the other packages, the implementation of the wide data format in [randomLCA](#) supports missing data. This feature is not implemented in [rater](#) because the long data format allows both missing data and repeated ratings.

Overall, the main features that distinguish [rater](#) from the other packages are its full support for the Dawid–Skene model and extensions, and its support for repeated ratings.

5 Implementation details

[rater](#) uses Stan (Carpenter et al. 2017) to fit the above models to data. It therefore supports both optimisation and Markov chain Monte Carlo (MCMC) sampling, in particular using the No U-Turn Sampler (NUTS) algorithm (Hoffman and Gelman 2014). For most datasets we recommend using NUTS, due to the fact that it will generate realisations of the full posterior distribution. Optimisation may be useful, however, for particularly large datasets, especially if K is large, or if there is a need to fit a large number of models.

5.1 Marginalisation

The NUTS algorithm relies on computing derivatives of the (log) posterior distribution with respect to all parameters. It cannot, therefore, be used when the posterior contains discrete parameters, such as the latent class in the Dawid–Skene model. To overcome this difficulty the Stan models implemented in [rater](#) use marginalised forms of the posterior to allow NUTS to be used. In other words, for the Dawid–Skene model we implement the likelihood $\Pr(y \mid \theta, \pi)$, as described earlier in [Section 4.1](#), together with priors for θ and π .

To avoid any confusion, we stress the fact that our choice to marginalise over the vector z is purely because it is discrete and we want to use the NUTS algorithm. It is not due to the fact that the components of the vector z are inherently latent variables. Alternative Bayesian implementations could avoid marginalisation and sample the z_i 's directly (although this may be less efficient, see [Section 6.3](#) for discussion). Similarly, if the components of z were all continuous, we would be able to sample them using NUTS. If one or more of the other non-latent parameters were discrete, then we would need to marginalise over them too.

5.2 Inference for the true class via conditioning

Marginalising over z allows us to use NUTS and sample efficiently from the posterior probability distribution of θ and π . We are still able to obtain the posterior distribution of the components z_i of z as follows. For each posterior draw of θ and π , we calculate the conditional posterior for z_i , $\Pr(z_i | \theta^*, \pi^*, y)$, where the conditioning is on the **drawn** values, θ^* and π^* . This can be done using Bayes' theorem, for example in the Dawid–Skene model,

$$\begin{aligned} \Pr(z_i = k | \theta^*, \pi^*, y) &= \frac{\Pr(y | z_i = k, \theta^*, \pi^*) \Pr(z_i = k | \theta^*, \pi^*)}{\Pr(y | \theta^*, \pi^*)} \\ &= \frac{\pi_k^* \prod_{j=1}^I \theta_{j,k,y_{i,j}}^*}{\sum_{m=1}^K \pi_m^* \prod_{j=1}^I \theta_{j,m,y_{i,j}}^*}. \end{aligned}$$

To get a Monte Carlo estimate of the marginal posterior, $\Pr(z_i | y)$, we take the mean of the above conditional probabilities across the posteriors draws of θ and π . To formally justify this step we can use the same argument that justifies the general MCMC technique of Rao–Blackwellization, see [Section 6](#). In practice, this calculation is straightforward with Stan because we can write the log posterior in terms of the (log scale) unnormalised versions of the conditional probabilities and output them alongside the posterior draws. Then it is simply a matter of renormalising them, followed by taking the mean across draws.

5.3 Data summarisation

[rater](#) implements the ability to use different forms of the likelihood depending on the format of the data that is available. For example, for a balanced design the full likelihood for the Dawid–Skene model is

$$\Pr(y | \theta, \pi) = \prod_{i=1}^I \left(\sum_{k=1}^K \left(\pi_k \cdot \prod_{j=1}^I \theta_{j,k,y_{i,j}} \right) \right).$$

Looking closely, we see that the contribution of any two items with the same pattern of ratings, that is, the same raters giving the same ratings, is identical. We can therefore rewrite it as a product over patterns rather than items,

$$\Pr(y | \theta, \pi) = \prod_{l=1}^L \left(\sum_{k=1}^K \left(\pi_k \cdot \prod_{j=1}^I \theta_{j,k,y_{l,j}} \right) \right)^{n_l},$$

where L is the number of distinct rating ‘patterns’, n_l is the number of times pattern l occurs and $y_{l,j}$ denotes the rating given by rater j within pattern l . This corresponds directly to having the data in the ‘grouped’ format, see [Section 2](#).

This rewriting is useful for balanced designs where there are many ratings but few distinct patterns, as it removes the need to iterate over every rating during computation of the posterior, reducing the time for model fitting. An example of this type of data are the caries data presented in [Section 7](#).

This technique is not always helpful. For example, if the data has many missing entries, the number of distinct patterns L may be similar or equal to than the total number of ratings. Consequently, rewriting to use the grouped format may not save much computation. For this reason, the implementation of [rater](#) does not allow missing values in grouped data. Note that in the long data format we simply drop rows that correspond to missing data, a result of which is that missing data is never explicitly represented in [rater](#).

Both the [randomLCA](#) and [BayesLCA](#) packages support grouped format data input. Of these, CRANpkg[randomLCA] supports missing values in grouped format data.

6 Marginalisation and Rao–Blackwellization

The technique of marginalisation for sampling discrete random variables discussed in [Section 5](#) is well known in the Stan community. We feel, however, that the interpretation and use of the relevant conditional probabilities and expectations has not been clearly described before, so we attempt to do so here. We begin with a brief review of the theory of marginalisation, which is a special case of a more general technique often referred to as ‘Rao–Blackwellization’. See Owen (2013) for an introduction and Robert and Roberts (2021) for a recent survey of the use of Rao–Blackwellization in MCMC.

6.1 Connection with the Rao–Blackwell theorem

Suppose we are interested in estimating $\mu = \mathbb{E}(f(X, Y))$ for some function f of random variables X and Y . An example of such an expectation is $\mu = \mathbb{E}(Y)$ where we take $f(x, y) = y$. If we can sample from the joint distribution of X and Y , $p_{X,Y}(x, y)$, the obvious Monte Carlo estimator of this expectation is

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n f(x_i, y_i)$$

where (x_i, y_i) are samples from the joint distribution. Now let $h(x) = \mathbb{E}(f(X, Y) \mid X = x)$. An alternate conditional estimator—the so-called *Rao–Blackwellized estimator*—of μ is

$$\hat{\mu}_{\text{cond}} = \frac{1}{n} \sum_{i=1}^n h(x_i)$$

where the x_i are sampled from the distribution of X . This estimator is justified by the fact that,

$$\mathbb{E}(h(X)) = \mathbb{E}(\mathbb{E}(f(X, Y) \mid X)) = \mathbb{E}(f(X, Y)) = \mu$$

For this to be effective we need to be able to efficiently compute (ideally, analytically) the conditional expectation $h(X)$. The original motivation for using this technique was to reduce the variance of Monte Carlo estimators. This follows from noting that

$$\begin{aligned} \text{var}(f(X, Y)) &= \mathbb{E}(\text{var}(f(X, Y) \mid X)) + \text{var}(\mathbb{E}(f(X, Y) \mid X)) \\ &= \mathbb{E}(\text{var}(f(X, Y) \mid X)) + \text{var}(h(X)), \end{aligned}$$

and since the first term on the right-hand side is non-negative, we have $\text{var}(f(X, Y)) \geq \text{var}(h(X))$.

From this result it follows that, if the draws x_i are independent samples, then

$$\text{var}(\hat{\mu}_{\text{cond}}) < \text{var}(\hat{\mu})$$

for all functions $f(X, Y)$. In the applications we consider here, however, x_i will be drawn using MCMC and therefore draws will not be independent. In this case it is possible to construct models and functions $f(\cdot)$ for which conditioning can increase the variance of the estimator (Geyer 1995).

The name ‘Rao–Blackwellization’ arose due to the similarity of the above argument to that used in the proof of the Rao–Blackwell theorem (Blackwell 1947), which states that given an estimator $\hat{\mu}$ of a parameter μ , the conditional expectation of $\hat{\mu}$ given a sufficient statistic for μ is potentially more efficient and certainly no less efficient than $\hat{\mu}$ as an estimator of μ . The Rao–Blackwellized estimators presented above do not, however, make any use of sufficiency and do not have the same optimality guarantees that the Rao–Blackwell theorem provides, making the name less than apt. Following Geyer (1995), we prefer to think of the estimators presented above as simply being formed from averaging conditional expectations.

6.2 Marginalisation in Stan

The motivation for marginalisation in Stan is to enable estimation of $\mu = \mathbb{E}(f(X, Y))$ without having to sample from (X, Y) if either X or Y is discrete. Suppose that Y is discrete and X continuous. To compute a Monte Carlo estimate of $\mathbb{E}(f(X, Y))$ using Stan we carry out four steps.

First, marginalise out Y from $p_{X,Y}(x, y)$ to give $p_X(x)$. (See Equation (2.4.1) in [Section 4.1](#) for how this is done in the Dawid–Skene model.) This marginal distribution, which only contains continuous parameters, should then be implemented as a Stan model. Second, using Stan as usual, sample from the distribution of the continuous parameters X to give Monte Carlo samples $\{x_i\}$. Given only the samples from the distribution of X , we can estimate $\mathbb{E}(f(X, Y))$ using the Rao–Blackwellized estimator described in the previous section. Doing so requires us to evaluate $h(X) = \mathbb{E}(f(X, Y) \mid X)$ for all

functions $f(\cdot)$ in which we may be interested; this is the third step of the process. Finally, as the fourth step, we can evaluate $h(x)$ for each of the Monte Carlo draws x_i and estimate $\mathbb{E}(f(X, Y))$ by

$$\frac{1}{n} \sum_{i=1}^n h(x_i).$$

Evaluating the conditional expectation

The third step outlined above may appear to require substantial mathematical manipulation. In practice, however, we can use the discrete nature of the latent class to simplify the calculation. Specifically, for any function $f(x, y)$ we have

$$h(X) = \mathbb{E}(f(X, Y) \mid X) = \sum_k f(X, k) \Pr(Y = k \mid X),$$

where we sum over the values in the latent discrete parameters. An essential part of this formula is the probability of the discrete parameters conditional on the continuous parameters, $\Pr(Y = k \mid X)$. This quantity can be derived easily through Bayes' theorem or can be encoded as part of the marginalised Stan model; see [Section 5.2](#) or the next section for how this is done in the case of the Dawid–Skene model.

In the Dawid–Skene model, and many other models with discrete variables, the discrete variables only take values in a finite set. Some models however may have discrete parameters which take countably infinite values, such as Poisson distributions. In this case, marginalisation is still possible but will generally involve approximating infinite sums, both when marginalising out the discrete random variables and calculating the expectations (as in the equation above), which may be computationally expensive.

Furthermore, for the cases where $f(x, y)$ is a function of only x or y , the general formula simplifies further. Firstly, when $f(x, y) = f(x)$ we have

$$h(X) = \sum_k f(X) \Pr(Y = k \mid X) = f(X) \sum_k \Pr(Y = k \mid X) = f(X).$$

This means that we can estimate $\mathbb{E}(f(X))$ with the standard, seemingly unconditional, estimator:

$$\sum_{i=1}^n f(x_i).$$

Even after marginalisation, computing expectations of functions of the continuous parameters can be performed as if no marginalisation had taken place.

Secondly, when $f(x, y) = f(y)$ we have that

$$h(X) = \sum_k f(k) \Pr(Y = k \mid X).$$

An important special case of this result is when $f(x, y) = \mathbf{1}(y = k)$, where $\mathbf{1}$ is the indicator function. This is important because it allows us to recover the probability mass function of the discrete random variable Y , since $\mathbb{E}(f(X, Y)) = \mathbb{E}(\mathbf{1}(Y = k)) = \Pr(Y = k)$. In this case we have

$$h(X) = \sum_k \mathbf{1}(y = k) \Pr(Y = k \mid X) = \Pr(Y = k \mid X).$$

We therefore estimate $\Pr(Y = k)$ with:

$$\frac{1}{n} \sum_{i=1}^n \Pr(Y = k \mid X = x_i).$$

Again, we stress that our ability to do these calculations relies upon being able to easily compute $\Pr(Y = k \mid X = x_i)$ for each of the Monte Carlo draws x_i .

Estimating the conditional probability of the true class

For the categorical rating problem (using the Dawid–Skene model), the discrete random variable of interest for each item i is the true class of the item, z_i . We use the technique from the previous section to calculate the posterior probability of the true class, as described in [Section 5.2](#). In this case, the discrete variable is z_i (taking on the role that Y played in the previous section), the continuous variables are

θ and π (which together take on the role that X played in the previous section), and all probability calculations are always conditional on the data (the ratings, y).

6.3 Efficiency of marginalisation

It is not immediately obvious whether marginalisation is more or less efficient than techniques that actually realise a value for the discrete random variable at each iteration, such as Gibbs sampling. Marginalising can be viewed as a form of Rao–Blackwellization, a general MCMC technique designed to reduce the variability of Monte Carlo estimators. This link strongly suggests that marginalisation is more efficient than using discrete random variables. Unfortunately, limitations of the theory of Rao–Blackwellization (see [Section 6.1](#)) and the difficulty of theoretically comparing different sampling algorithms, such as Gibbs sampling and NUTS, means that it is unclear whether marginalisation will always be computationally superior for a given problem.

However, in practice marginalisation does seem to improve convergence at least for non-conjugate models. For example, Yackulic et al. (2020) show that for the Cormack–Jolly–Seber model marginalisation greatly speeds up inference in JAGS, BUGS and WinBUGS. They also demonstrate that the marginalised model implemented in Stan is orders of magnitude more efficient than the marginalised models in the other languages. These results show that marginalisation has the potential to speed up classic algorithms and also allows the use of more efficient gradient-based algorithms, such as NUTS, for problems involving discrete random variables.

A recent similar exploration of marginalisation for some simple mixture models and the Dawid–Skene model, using JAGS and Stan, suggests that the software implementation has a greater impact on efficiency than the choice of whether or not marginalisation is used (Zhang et al. 2022). In their comparisons, Stan usually achieved the best performance (and note that Stan requires marginalisation of discrete parameters).

For these reasons, we recommend that practitioners using models with discrete parameters consider marginalisation if more efficiency is desired, and implement their models using Stan rather than JAGS.

7 Example usage

To demonstrate [rater](#) we use two example datasets.

The first dataset is taken from the original paper introducing the Dawid–Skene model (Dawid and Skene 1979). The data consist of ratings, on a 4-point scale, made by five anaesthetists of patients’ pre-operative health. The ratings were based on the anaesthetists assessments of a standard form completed for all of the patients. There are 45 patients (items) and four anaesthetists (raters) in total. The first anaesthetist assessed the forms a total of three times, spaced several weeks apart. The other anaesthetists each assessed the forms once. As in the Dawid–Skene paper, we will not seek to model the effect of time on the ratings of the first anaesthetist.

First we load the [rater](#) package:

```
library(rater)

#> * The rater package uses `Stan` to fit bayesian models.
#> * If you are working on a local, multicore CPU with excess RAM please call:
#> * options(mc.cores = parallel::detectCores())
#> * This will allow Stan to run inference on multiple cores in parallel.
```

This will display information about altering options used by Stan to make best use of computational resources on your machine. We can then load and look at a snippet of the anaesthesia data, which is included in the package.

```
data("anesthesia", package = "rater")
head(anesthesia)
```

```
#>   item rater rating
#> 1    1     1      1
#> 2    1     1      1
#> 3    1     1      1
#> 4    1     2      1
#> 5    1     3      1
#> 6    1     4      1
```

These data are arranged in ‘long’ format where each row corresponds to a single rating. The first column gives the index of the item that was rated, the second the index of the rater that rated the item and the third the actual rating that was given. This layout of the data supports raters rating the same item multiple times, as happens in the dataset. It is also the most convenient from the perspective of fitting models but may not always be the optimal way to store or represent categorical rating data; see [Using grouped data](#) for an example using the ‘grouped’ data format.

7.1 Fitting the model

We can fit the Dawid–Skene model using MCMC by running the command:

```
fit_1 <- rater(anesthesia, dawid_skene())
```

This command will print the running output from Stan, providing an indication of the progress of the sampler (for brevity, we have not shown this output here). To fit the model via optimisation, we set the method argument to “optim”:

```
fit_2 <- rater(anesthesia, dawid_skene(), method = "optim")
```

The second argument of the `rater()` function specifies the model to use. We have implemented this similarly to the `family` argument in `glm()`, which can be passed as either a function or as a character string. The above examples pass the model as a function. We could have instead passed it as a string, for example:

```
fit_2 <- rater(anesthesia, "dawid_skene", method = "optim")
```

Either version will fit the Dawid–Skene model using the default choice of prior. The benefit of passing the model as a function is that it allows you to change the prior, see [Different models and priors](#).

7.2 Inspecting the fitted model

`rater` includes several ways to inspect the output of fitted models. These are summarised in [Table 4](#), and we illustrate many of them here. Firstly, we can generate a text summary:

```
summary(fit_1)

#> Model:
#> Bayesian Dawid and Skene Model
#>
#> Prior parameters:
#>
#> alpha: default
#> beta: default
#>
#> Fitting method: MCMC
#>
#> pi/theta samples:
#>      mean   5%  95% Rhat ess_bulk
#> pi[1]    0.37 0.27 0.48    1  9094.91
#> pi[2]    0.41 0.30 0.52    1  8050.47
#> pi[3]    0.14 0.07 0.23    1  6864.65
#> pi[4]    0.07 0.03 0.14    1  6665.16
#> theta[1, 1, 1] 0.86 0.79 0.93    1  7728.01
#> theta[1, 1, 2] 0.10 0.05 0.17    1  7857.38
#> theta[1, 1, 3] 0.02 0.00 0.05    1  5619.81
#> theta[1, 1, 4] 0.02 0.00 0.05    1  6283.33
#> # ... with 76 more rows
#>
#> z:
#>      MAP Pr(z = 1) Pr(z = 2) Pr(z = 3) Pr(z = 4)
#> z[1]    1      1.00      0.00      0.00      0.00
```

Table 4: Methods to inspect fitted models and what values they return.

	MCMC mode	Optimisation mode
'summary()'	Basic information about the fitted model	Basic information about the fitted model
'point_estimate()'	Posterior means for π and θ , posterior modes for z	Posterior modes for z
'posterior_interval()'	Credible intervals for π and θ	N/A
'posterior_samples()'	MCMC draws for π and θ	N/A
'mcmc_diagnostics()'	MCMC convergence diagnostics for π and θ	N/A
'class_probabilities()'	Posterior distribution for z	Posterior distribution for z

```
#> z[2] 3 0.00 0.00 0.98 0.02
#> z[3] 2 0.38 0.62 0.00 0.00
#> z[4] 2 0.01 0.99 0.00 0.00
#> z[5] 2 0.00 1.00 0.00 0.00
#> z[6] 2 0.00 1.00 0.00 0.00
#> z[7] 1 1.00 0.00 0.00 0.00
#> z[8] 3 0.00 0.00 1.00 0.00
#> # ... with 37 more items
```

This function will show information about which model has been fitted and the values of the prior parameters that were used. The displayed text also contains information about the parameter estimates and posterior distributions, and the convergence of the sampler.

We can extract point estimates for any of the parameters or the latent classes via the `point_estimate()` function. For example, the following will return the latent class with the highest posterior probability (i.e., the posterior modes) for each item:

```
point_estimate(fit_1, "z")

#> $z
#> 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26
#> 1 3 2 2 2 2 1 3 2 2 4 2 1 2 1 1 1 1 2 2 2 2 2 1 1
#> 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45
#> 2 1 1 1 1 3 1 2 2 3 2 2 3 1 1 1 2 1 2
```

The following will return the posterior means of the prevalence probabilities:

```
point_estimate(fit_1, "pi")

#> $pi
#> [1] 0.37432122 0.40674116 0.14406295 0.07487466
```

From the outputs above, we can see that the model has inferred that most patients have health that should be classified as category 1 or 2 (roughly 40% each), and categories 3 and 4 being rarer (about 14% and 7% respectively). Based on the point estimates, there are examples of patients from each category in the dataset.

The function call `point_estimate(fit1, "theta")` will return posterior means for the parameters in the error matrices (not shown here, for brevity). When used with optimisation fits, `point_estimate()` will return posterior modes rather than posterior means for the parameters.

7.3 Inspecting posterior distributions

To represent uncertainty, we need to look beyond point estimates. The function `posterior_interval()` will return credible intervals for the parameters. For example, 80% credible intervals for the elements of the error matrices:

```
head(posterior_interval(fit_1, 0.8, "theta"))

#>
#>          10%          90%
#> theta[1, 1, 1] 0.805074415 0.91666962
#> theta[1, 1, 2] 0.055802572 0.15309851
```



```
#> theta[1, 1, 3] 0.001813371 0.03941595
#> theta[1, 1, 4] 0.001760845 0.03832660
#> theta[1, 2, 1] 0.026218408 0.10590753
#> theta[1, 2, 2] 0.791396057 0.90637035
```

The function `posterior_samples()` will return the actual MCMC draws. For example (here just illustrating the draws of the π parameter):

```
head(posterior_samples(fit_1, "pi")$pi)

#>
#> iterations      [,1]      [,2]      [,3]      [,4]
#>      [1,] 0.4830595 0.3579169 0.1254949 0.03352871
#>      [2,] 0.2939839 0.3834238 0.2516929 0.07089945
#>      [3,] 0.3636097 0.4148082 0.1309119 0.09067012
#>      [4,] 0.3798456 0.3133947 0.2353036 0.07145615
#>      [5,] 0.3639428 0.4291340 0.1760525 0.03087066
#>      [6,] 0.3845465 0.3831443 0.1721656 0.06014354
```

Neither of the above will work for optimisation fits, which are limited to point estimates only. For the latent classes, we can produce the probability distribution as follows:

```
head(class_probabilities(fit_1))

#>
#>      [,1]      [,2]      [,3]      [,4]
#> 1 9.999994e-01 1.340790e-07 2.916232e-08 4.521956e-07
#> 2 8.469190e-08 2.570243e-05 9.773942e-01 2.257999e-02
#> 3 3.823085e-01 6.171346e-01 1.075246e-04 4.493620e-04
#> 4 5.252130e-03 9.942742e-01 3.418545e-04 1.318565e-04
#> 5 2.473083e-07 9.999633e-01 3.338995e-05 3.062785e-06
#> 6 1.674456e-06 9.993819e-01 5.813534e-04 3.502596e-05
```

This works for both MCMC and optimisation fits. For the former, the output is the posterior distribution on z , while for the latter it is the distribution conditional on the point estimates of the parameters (π and θ).

7.4 Plots

It is often easier to interpret model outputs visually. Using [rater](#), we can plot the parameter estimates and distribution of latent classes.

The following command will visualise the posterior means of the error rate parameters, with the output shown in [Figure 2](#):

```
plot(fit_1, "raters")
```

We can see high values along the diagonals of these matrices, which indicates that each of the 5 anaesthetists is inferred as being fairly accurate at rating pre-operative health. Looking at the non-diagonal entries, we can see that typically the most common errors are 1-point differences on the rating scale.

The following command visualises the latent class probabilities, with the output shown in [Figure 3](#):

```
plot(fit_1, "latent_class", item_index = c(2, 3, 12, 36, 38))
```

For the purpose of illustration, for this plot we have selected the 5 patients with the greatest uncertainty in their pre-operative health (latent class). The other 40 patients all have almost no posterior uncertainty for their pre-operative health. Thus, suppose we wished to use the model to infer the pre-operative health by combining all of the anaesthetists' ratings, we can do so confidently for all but a handful of patients.

The same output for the models fitted via optimisation rather than MCMC (using `fit_2` instead of `fit_1`) are shown in [Figure 4](#) and [Figure 5](#). We can see that this estimation method leads to the same




Figure 2: Visual representation of the inferred parameters in the error matrices (θ) for the Dawid-Skene model fitted via MCMC to the anaesthesia dataset. The values shown are posterior means, with each cell shaded on a gradient from white (value close to 0) to dark blue (value close to 1).

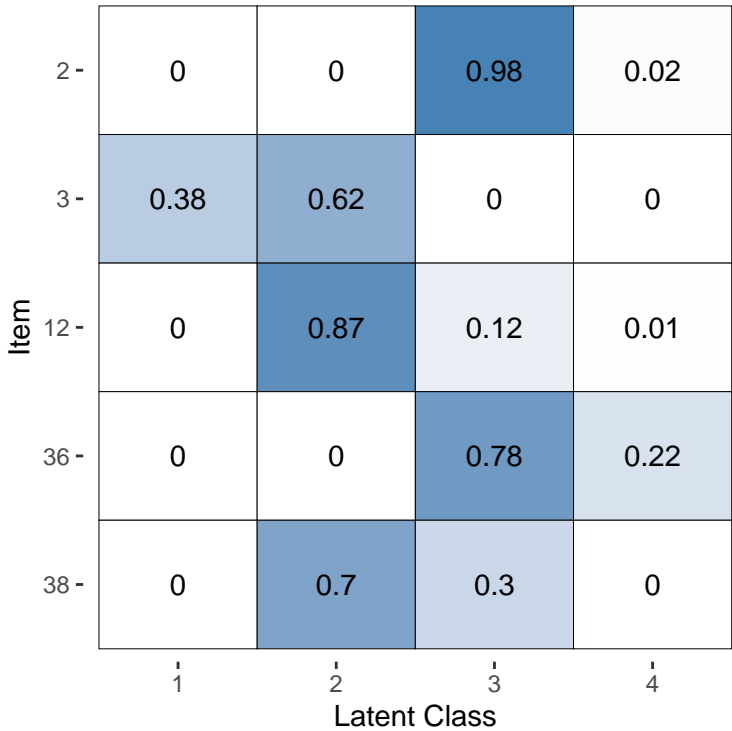


Figure 3: Visualisation of the inferred probability of each latent class, for a selected subset of items, for the Dawid-Skene model fitted via MCMC to the anaesthesia dataset.




Figure 4: Visual representation of the inferred parameters in the error matrices (θ) for the Dawid-Skene model fitted via optimisation to the anaesthesia dataset. Compare with [Figure 2](#), which used MCMC instead of optimisation.




Figure 5: Visualisation of the inferred probability of each latent class, for a selected subset of items, for the Dawid-Skene model fitted via optimisation to the anaesthesia dataset. Compare with [Figure 3](#), which used MCMC instead of optimisation.




Figure 6: Visualisation of the inferred population prevalence parameters, with 90% credible intervals, for the Dawid–Skene model fitted to the anaesthesia dataset.

broad conclusions, however the optimisation-based estimates have considerably less uncertainty: they are typically closer to 0 or 1. This behaviour reflects the fact that optimisation-based inference will usually not capture the full uncertainty and will lead to overconfident estimates. Thus, we recommend using MCMC (which we have set as the default).

While it is typically of less interest, it is also possible to visualise the prevalence estimates, together with credible intervals, using the following command (see output in [Figure 6](#)):

```
plot(fit_1, "prevalence")
```

7.5 Different models and priors

The second argument to the `rater()` function specifies what model to use, including details of the prior if not using the default one. This gives a unified place to specify both the model and prior.

For example, this is how to set a different prior using the Dawid–Skene model:

```
diff_alpha_fit <- rater(anaesthesia, dawid_skene(alpha = rep(10, 4)))
```

When specifying the β hyper-parameters for the Dawid–Skene model, the user can either specify a single matrix or a 3-dimensional array. If only a matrix specified, it will be interpreted as the hyper-parameter values for all raters. This is useful in the common situation where the overall quality of the raters is known but there is no information on the quality of any specific rater. When a 3-dimensional array is passed, it is taken to specify the hyper-parameters for each of the raters (i.e., a separate matrix for each rater). This is useful when prior information about the quality of specific raters is available, for example when some ‘raters’ are diagnostic tests with known performance characteristics.

This is how to use the class-conditional Dawid–Skene model (with default prior):

```
diff_model_fit <- rater(anaesthesia, class_conditional_dawid_skene())
```

Compared with the Dawid–Skene model, the latter uses error matrices with the constraint that all off-diagonal entries must be the same. We can visualise the θ parameter using the following command (output in [Figure 7](#)) and see that all off-diagonal elements are indeed equal:

```
plot(diff_model_fit, "theta")
```

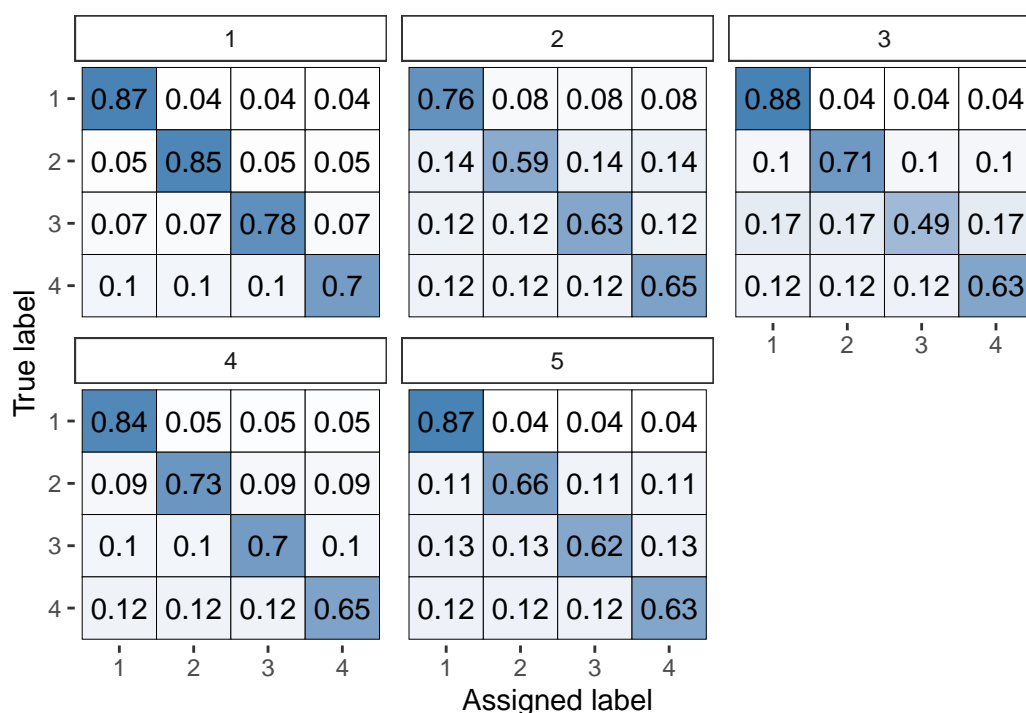


Figure 7: Visual representation of the inferred parameters in the error matrices (θ) for the class-conditional Dawid–Skene model fitted via MCMC to the anaesthesia dataset. Compare with [Figure 2](#), which used the standard Dawid–Skene model.

7.6 Using grouped data

The second example dataset is taken from Espeland and Handelman (1989). It consists of 3,689 binary ratings, made by 5 dentists, of whether a given tooth was healthy or had caries/cavities. The ratings were performed using X-ray only, which was thought to be more error-prone than visual/tactile assessment of each tooth (see Handelman et al. (1986) for more information and a description of the wider dataset from which these binary ratings were taken).

```
data("caries", package = "rater")
head(caries)
```

```
#>   rater_1 rater_2 rater_3 rater_4 rater_5    n
#> 1      1      1      1      1      1 1880
#> 2      1      1      1      1      2  789
#> 3      1      1      1      2      1   43
#> 4      1      1      1      2      2   75
#> 5      1      1      2      1      1   23
#> 6      1      1      2      1      2   63
```

This is an example of ‘grouped’ data. Each row represents a particular ratings ‘pattern’, with the final column being a tally that records how many instances of that pattern appear in the data. **rater** accepts data in either ‘long’, ‘wide’ or ‘grouped’ format. The ‘long’ format is the default because it can represent data with repeated ratings. When available the ‘grouped’ format can greatly speed up the computation of certain models and is convenient for datasets that are already recorded in that format. The ‘wide’ format doesn’t provide any computational advantages but is a common format for data without repeated ratings and so is provided for convenience.

Here’s how we fit the model to grouped data (note the `data_format` argument):

```
fit3 <- rater(caries, dawid_skene(), data_format = "grouped")
```

From there, we can inspect the model fit in the same way as before. For example (output shown in [Figure 8](#)):

```
plot(fit3, "raters")
```




Figure 8: Visual representation of the inferred parameters in the error matrices (θ) for the Dawid–Skene model fitted via MCMC to the caries dataset.

The first 4 dentists are highly accurate at diagnosing healthy teeth (rating = 1), but less so at diagnosing caries (rating = 2). In contrast, the 5th dentist was much better at diagnosing caries than healthy teeth.

7.7 Convergence diagnostics

A key part of applied Bayesian statistics is assessing whether the MCMC sampler has converged on the posterior distribution. To summarise the convergence of a model fit using MCMC, `rater` provides the `mcmc_diagnostics()` function.

```
head(mcmc_diagnostics(fit_1))
```

```
#>               Rhat ess_bulk
#> pi[1]          0.9997012 9094.906
#> pi[2]          1.0009095 8050.465
#> pi[3]          1.0000406 6864.648
#> pi[4]          1.0021824 6665.159
#> theta[1, 1, 1] 1.0013982 7728.007
#> theta[1, 1, 2] 1.0011678 7857.378
```

This function calculates and displays the \hat{R} statistic (Rhat) and bulk effective sample size (ess_bulk) for all of the π and θ parameters. Users can then check for the convergence of specific parameters by applying standard rules, such as considering that convergence has been reached for a specific parameter if its $\hat{R} < 1.01$ (Vehtari et al. 2021).

Users wishing to calculate other metrics, or produce convergence visualisations such as trace plots, can either extract the underlying Stan model from the `rater` fit using the `get_stanmodel()` function, or convert the fit into a `mcmc.list` from the `coda` package using the `as_mcmc.list()` function. Functions in the `rstan` and `coda` packages will then allow visualisation based assessment of convergence and the calculation of other, more advanced, diagnostics.

For example, the following code uses `rstan` to draw the trace plot (shown in Figure 9) for one of the parameters from the Dawid–Skene model fitted to the anaesthesia data. We can see that the four chains show good mixing. The trace plots for the other parameters (not shown here) are similar, indicating that the MCMC sampling procedure has converged.




Figure 9: A trace plot for the π_1 parameter from the Dawid-Skene model fitted via MCMC to the anaesthesia dataset.

```
stan_fit <- get_stanfit(fit_1)
rstan::traceplot(stan_fit, pars = "pi[1]")
```

7.8 Model assessment and comparison

Finally, [rater](#) provides facilities to assess and compare fitted models. First, it provides the `posterior_predict()` function to simulate from the posterior predictive distributions of all the models implemented in [rater](#). The simulated data can then be used, for example, to perform posterior predictive checks. These checks compare the data simulated from the fitted model to the observed data. A number of datasets are simulated from the posterior predictive distribution and a summary statistic is calculated for each dataset. The same summary statistic is calculated on the observed dataset and is compared to the distribution of simulated statistics. A large discrepancy between the simulated statistics and the observed statistic indicates possible model misspecification.

To illustrate this functionality, we consider assessing the fit of the Dawid-Skene model to the anaesthesia dataset. In this example, we simulate 1,000 datasets from the posterior predictive distribution, and use the proportion of the ratings that are class 2 as a statistic to summarise each dataset.

```
class2prop <- function() {
  simdata <- posterior_predict(fit_1, anaesthesia[, 1:2])
  sum(simdata$rating == 2) / nrow(simdata)
}
ppc_statistics <- replicate(1000, class2prop())
head(ppc_statistics)

#> [1] 0.5587302 0.4507937 0.4126984 0.4190476 0.4666667 0.2539683
```

We can then graphically compare the distribution of these statistics to the value of the same statistic applied to the anaesthesia dataset. The following commands will create such a plot (shown in [Figure 10](#)):

```
ggplot(data.frame(prop_class_two = ppc_statistics), aes(prop_class_two)) +
  geom_histogram(binwidth = 0.01, fill = "steelblue", colour = "black") +
  geom_vline(xintercept = sum(anaesthesia$rating == 2) / nrow(anaesthesia),
    colour = "black", linewidth = 2) +
  theme_bw() +
  coord_cartesian(xlim = c(0.1, 0.6)) +
  labs(
    x = "Proportion of class 2 ratings",
    y = "Count"
  )
```




Figure 10: An example of a posterior predictive check for the Dawid–Skene model fitted via MCMC to the anaesthesia dataset. Shown is a histogram of 1,000 simulated datasets from the posterior predictive distribution, each summarised by the proportion of ratings that are class 2. The vertical black line shows the proportion of ratings that are class 2 in the original dataset, for comparison against the histogram.

The statistic calculated on the anaesthesia dataset lies towards the centre of the distribution of statistics calculated from datasets drawn from the posterior predictive distribution. This is consistent with the model fitting the data well. (This is an illustrative example only. A more comprehensive set of such comparisons should be conducted before concluding that the model adequately describes the data.)

Second, `rater` implements the functions `loo()` and `waic()` to provide model comparison metrics using the `loo` package. The function `loo()` calculates an approximation to leave-out-one cross-validation, a measure of model performance using Pareto Smoothed Importance Sampling (Vehtari, Gelman, and Gabry 2017). The function `waic()` calculates measures related to the Widely Applicable Information Criterion (Watanabe 2013). Because `rater` fits models in a Bayesian framework, information criteria such as the AIC and BIC are not implemented.

In the context of `rater`, model comparison is only possible between the different implemented models, not between the same model with different parameter values or included covariates, as is more commonly the case. For this reason, considerations of data size and what is known about the characteristics of the raters should also be taken into account when making such choices, in addition to the output of `loo()` and `waic()`.

To briefly illustrate some of this functionality, we compare the standard Dawid–Skene model to the class-conditional one.

```
loo(fit_1)

#>
#> Computed from 4000 by 45 log-likelihood matrix
#>
#>      Estimate   SE
#> elpd_loo  -233.7 16.9
#> p_loo      19.8  2.6
#> looic      467.5 33.9
#> -----
#> Monte Carlo SE of elpd_loo is 0.1.
#>
#> Pareto k diagnostic values:
```

```

#>
#> (-Inf, 0.5] (good) 43 95.6% 734
#> (0.5, 0.7] (ok) 2 4.4% 497
#> (0.7, 1] (bad) 0 0.0% <NA>
#> (1, Inf) (very bad) 0 0.0% <NA>
#>
#> All Pareto k estimates are ok (k < 0.7).
#> See help('pareto-k-diagnostic') for details.

loo(diff_model_fit)

#>
#> Computed from 4000 by 45 log-likelihood matrix
#>
#> Estimate SE
#> elpd_loo -245.7 18.1
#> p_loo 10.3 1.2
#> looic 491.3 36.2
#> -----
#> Monte Carlo SE of elpd_loo is 0.1.
#>
#> All Pareto k estimates are good (k < 0.5).
#> See help('pareto-k-diagnostic') for details.

loo_compare(loo(fit_1), loo(diff_model_fit))

#> elpd_diff se_diff
#> model1 0.0 0.0
#> model2 -11.9 3.2

```

From these results, we see that both models fit the data well, with the class-conditional model being slightly preferred.

8 Summary and discussion

Rating procedures, in which items are sorted into categories, are subject to both classification error and uncertainty when the categories themselves are defined subjectively. Data that results from these types of tasks require a proper statistical treatment if: (1) the population frequencies of the categories are to be estimated with low bias, (2) items are to be assigned to these categories reliably, and (3) the agreement of raters is to be assessed. To date there have been few options for practitioners seeking software that implements the required range of statistical models. The R package described in this paper, [rater](#), provides the ability to fit Bayesian versions of a variety of statistical models to categorical rating data, with a range of tools to extract and visualise parameter estimates.

The statistical models that we have presented are based on the Dawid–Skene model and recent modifications of it including extensions, such as the hierarchical Dawid–Skene model where the rater-specific parameters are assumed to be drawn from a population distribution, and simplifications that, for example, assume exchangeable raters with identical performance characteristics (homogeneous Dawid–Skene), or homogeneity criteria where classification probabilities are identical regardless of the true, underlying category of an item (class-conditional Dawid–Skene).

We provided: (1) an explanation of the type of data formats that categorical ratings are recorded in, (2) a description of the construction and implementation of the package, (3) a comparison of our package to other existing packages for fitting latent class models in R, (4) an introduction to the user interface, and (5) worked examples of real-world data analysis using [rater](#).

We devoted an entire section to motivating, deriving and explaining the use of a marginalised version of the joint posterior probability distribution to remove dependence on the unknown value of the true, underlying rating category of an item. This is necessary because the No-U-Turn Sampler, the main MCMC algorithm used in Stan, relies on computing derivatives with respect to all parameters that must, therefore, be continuously-valued. The technique involves the use of conditional expectation and is a special case of a more general technique of conditioning or ‘Rao–Blackwellization’, the process of transforming an estimator using the Rao–Blackwell theorem to improve its efficiency.

Our package was developed with the classification of medical images in mind, where there may be a large number of images but typically only a small number of possible categories and a limited

number of expert raters. The techniques proposed based on extensions of the Dawid–Skene model readily extend to scenarios where the datasets are much larger, or where the raters behave erratically (cannot be relied on to rate correctly more frequently than they rate incorrectly) or are antagonistic (deliberately attempt to allocate items to a category not favoured by the majority of raters).

One possible extension of the Dawid–Skene model is to add item- and rater-specific covariates. These covariates would encode the difficulty of items and the expertise (or lack thereof) of the raters. This extension is particularly attractive as it would partially alleviate the strong assumption of independence conditional on the latent class, replacing it with the weaker assumption of independence conditional on the latent class and the covariates. Unfortunately, these types of models would contain many more parameters than the original Dawid–Skene model making them difficult to fit, especially when used with relatively small datasets common in the context of rating data in medical research. Therefore, future methodological research is needed before these models can be included in the **rater** package. Latent class models with covariates can be fitted, although only in a frequentist framework, using the R package **poLCA**.

9 Computational details

The results in this paper were obtained using R 4.2.2 (R Core Team 2021) and the following packages: **coda** 0.19.4 (Plummer et al. 2006), **ggplot2** 3.4.3 (Wickham 2016), **knitr** 1.45 (Xie 2014, 2015, 2023), **loo** 2.6.0 (Vehtari et al. 2023), **rater** 1.3.1 (Pullin and Vukcevic 2023), **rjtools** 1.0.12 (O’Hara-Wild et al. 2023), **rmarkdown** 1.0.12 (Xie, Allaire, and Golemund 2018; Xie, Dervieux, and Riederer 2020; Allaire et al. 2023), **rstan** 2.26.23 (Stan Development Team 2023). R itself and all packages used are available from the **Comprehensive R Archive Network** (CRAN).

10 Acknowledgments

We would like to thank Bob Carpenter for many helpful suggestions, and David Whitelaw for his flexibility in allowing the first author to work on this paper while employed at the Australian Institute of Heath and Welfare. Thanks also to Lars Mølgaard Saxhaug for his code contributions to **rater**.

References

- Allaire, JJ, Yihui Xie, Christophe Dervieux, Jonathan McPherson, Javier Luraschi, Kevin Ushey, Aron Atkins, et al. 2023. *rmarkdown: Dynamic Documents for r*. <https://github.com/rstudio/rmarkdown>.
- Asselineau, J., A. Paye, E. Bessède, P. Perez, and C. Proust-Lima. 2018. “Different Latent Class Models Were Used and Evaluated for Assessing the Accuracy of *Campylobacter* Diagnostic Tests: Overcoming Imperfect Reference Standards?” *Epidemiology and Infection* 146 (12): 1556–64. <https://doi.org/10.1017/S0950268818001723>.
- Beath, Ken J. 2017. “randomLCA: An R Package for Latent Class with Random Effects Analysis.” *Journal of Statistical Software* 81 (13): 1–25. <https://doi.org/10.18637/jss.v081.i13>.
- Berkes, Pietro, Bob Carpenter, Andrey Rzhetsky, and James Evans. 2011. <http://docs.enthought.com/uchicago-pyanno/>.
- Blackwell, David. 1947. “Conditional Expectation and Unbiased Sequential Estimation.” *Ann. Math. Statist.* 18 (1): 105–10. <https://doi.org/10.1214/aoms/1177730497>.
- Carpenter, Bob, Andrew Gelman, Matthew D. Hoffman, Daniel Lee, Ben Goodrich, Michael Betancourt, Marcus Brubaker, Jiqiang Guo, Peter Li, and Allen Riddell. 2017. “Stan: A Probabilistic Programming Language.” *Journal of Statistical Software* 76 (1). <https://doi.org/10.18637/jss.v076.i01>.
- Dawid, A. P., and A. M. Skene. 1979. “Maximum Likelihood Estimation of Observer Error-Rates Using the EM Algorithm.” *Applied Statistics* 28 (1): 20. <https://doi.org/10.2307/2346806>.
- Espeland, Mark A., and Stanley L. Handelman. 1989. “Using Latent Class Models to Characterize and Assess Relative Error in Discrete Measurements.” *Biometrics* 45 (2): 587–99. <https://doi.org/10.2307/2531499>.
- Geyer, Charles J. 1995. “Conditioning in Markov Chain Monte Carlo.” *Journal of Computational and Graphical Statistics* 4 (2): 148–54. <https://doi.org/10.1080/10618600.1995.10474672>.
- Goodman, Leo A. 1974. “Exploratory Latent Structure Analysis Using Both Identifiable and Unidentifiable Models.” *Biometrika* 61 (2): 215–31. <https://doi.org/10.1093/biomet/61.2.215>.
- Handelman, Stanley L., D. H. Leverett, Mark A. Espeland, and Jennifer A. Curzon. 1986. “Clinical Radiographic Evaluation of Sealed Carious and Sound Tooth Surfaces.” *The Journal of the American Dental Association* 113 (5): 751–54. <https://doi.org/10.14219/jada.archive.1986.0269>.

- Hoffman, Matthew D., and Andrew Gelman. 2014. "The No-U-Turn Sampler: Adaptively Setting Path Lengths in Hamiltonian Monte Carlo." *Journal of Machine Learning Research* 15: 1593–623. <http://jmlr.org/papers/v15/hoffman14a.html>.
- Hui, S. L., and S. D. Walter. 1980. "Estimating the Error Rates of Diagnostic Tests." *Biometrics* 36 (1): 167–71. <https://doi.org/10.2307/2530508>.
- Ipeirotis, Panagiotis G, Foster Provost, and Jing Wang. 2010. "Quality Management on Amazon Mechanical Turk." In *Proceedings of the ACM SIGKDD Workshop on Human Computation*, 64–67.
- Jakobsdottir, Johanna, and Daniel E Weeks. 2007. "Estimating Prevalence, False-Positive Rate, and False-Negative Rate with Use of Repeated Testing When True Responses Are Unknown." *The American Journal of Human Genetics* 81 (5): 1111–13. <https://doi.org/10.1086/521582>.
- Li, Hongwei, and Bin Yu. 2014. "Error Rate Bounds and Iterative Weighted Majority Voting for Crowdsourcing." *arXiv Preprint arXiv:1411.4086*. <https://arxiv.org/abs/1411.4086>.
- Linzer, Drew A., and Jeffrey B. Lewis. 2011. "poLCA: An R Package for Polytomous Variable Latent Class Analysis." *Journal of Statistical Software* 42 (1): 1–29. <https://doi.org/10.18637/jss.v042.i10>.
- O'Hara-Wild, Mitchell, Stephanie Kobakian, H. Sherry Zhang, Di Cook, Simon Urbanek, and Christophe Dervieux. 2023. *rjtools: Preparing, Checking, and Submitting Articles to the "R Journal"*. <https://CRAN.R-project.org/package=rjtools>.
- Owen, Art B. 2013. *Monte Carlo Theory, Methods and Examples*. <https://statweb.stanford.edu/~owen/mc/>.
- Passonneau, Rebecca J., and Bob Carpenter. 2014. "The Benefits of a Model of Annotation." *Transactions of the Association for Computational Linguistics* 2 (December): 311–26. https://doi.org/10.1162/tacl_a_00185.
- Paun, Silviu, Bob Carpenter, Jon Chamberlain, Dirk Hovy, Udo Kruschwitz, and Massimo Poesio. 2018. "Comparing Bayesian Models of Annotation." *Transactions of the Association for Computational Linguistics* 6 (December): 571–85. https://doi.org/10.1162/tacl_a_00040.
- Plummer, Martyn, Nicky Best, Kate Cowles, and Karen Vines. 2006. "CODA: Convergence Diagnosis and Output Analysis for MCMC." *R News* 6 (1): 7–11. <https://journal.r-project.org/archive/>.
- Pullin, Jeffrey, and Damjan Vukcevic. 2023. *rater: Statistical Models of Repeated Categorical Rating Data*. <https://CRAN.R-project.org/package=rater>.
- R Core Team. 2021. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.
- Robert, Christian P., and Gareth O. Roberts. 2021. "Rao-Blackwellization in the MCMC Era." *arXiv:2101.01011 [Stat]*. <https://arxiv.org/abs/2101.01011>.
- Rzhetsky, Hagit AND Wilbur, Andrey AND Shatkay. 2009. "How to Get the Most Out of Your Curation Effort." *PLOS Computational Biology* 5 (5): 1–13. <https://doi.org/10.1371/journal.pcbi.1000391>.
- Smyth, Padhraic, Usama Fayyad, Michael Burl, Pietro Perona, and Pierre Baldi. 1994. "Inferring Ground Truth from Subjective Labelling of Venus Images." In *Advances in Neural Information Processing Systems*, edited by G. Tesauro, D. Touretzky, and T. Leen. Vol. 7. MIT Press. https://proceedings.neurips.cc/paper_files/paper/1994/file/3cef96dcc9b8035d23f69e30bb19218a-Paper.pdf.
- Stan Development Team. 2021. *Stan User's Guide*. https://mc-stan.org/docs/2_27/stan-users-guide/.
- . 2023. "RStan: The R Interface to Stan." <http://mc-stan.org/>.
- Vehtari, Aki, Jonah Gabry, Mans Magnusson, Yuling Yao, Paul-Christian Bürkner, Topi Paananen, and Andrew Gelman. 2023. "Loo: Efficient Leave-One-Out Cross-Validation and WAIC for Bayesian Models." <https://mc-stan.org/loo/>.
- Vehtari, Aki, Andrew Gelman, and Jonah Gabry. 2017. "Practical Bayesian Model Evaluation Using Leave-One-Out Cross-Validation and WAIC." *Statistics and Computing* 27 (5): 1413–32. <https://doi.org/10.1007/s11222-016-9696-4>.
- Vehtari, Aki, Andrew Gelman, Daniel Simpson, Bob Carpenter, and Paul-Christian Bürkner. 2021. "Rank-Normalization, Folding, and Localization: An Improved \hat{R}^2 for Assessing Convergence of MCMC (with Discussion)." *Bayesian Analysis* 16 (2): 667–718. <https://doi.org/10.1214/20-BA1221>.
- Watanabe, Sumio. 2013. "A Widely Applicable Bayesian Information Criterion." *Journal of Machine Learning Research* 14 (March): 867–97.
- White, Arthur, and Thomas Brendan Murphy. 2014. "BayesLCA: An R Package for Bayesian Latent Class Analysis." *Journal of Statistical Software* 61 (13). <https://doi.org/10.18637/jss.v061.i13>.
- Wickham, Hadley. 2016. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York. <https://ggplot2.tidyverse.org>.
- Xie, Yihui. 2014. "knitr: A Comprehensive Tool for Reproducible Research in R." In *Implementing Reproducible Computational Research*, edited by Victoria Stodden, Friedrich Leisch, and Roger D. Peng. Chapman; Hall/CRC.

- . 2015. *Dynamic Documents with R and Knitr*. 2nd ed. Boca Raton, Florida: Chapman; Hall/CRC. <https://yihui.org/knitr/>.
- . 2023. *knitr: A General-Purpose Package for Dynamic Report Generation in r*. <https://yihui.org/knitr/>.
- Xie, Yihui, J. J. Allaire, and Garrett Grolemund. 2018. *R Markdown: The Definitive Guide*. Boca Raton, Florida: Chapman; Hall/CRC. <https://bookdown.org/yihui/rmarkdown>.
- Xie, Yihui, Christophe Dervieux, and Emily Riederer. 2020. *R Markdown Cookbook*. Boca Raton, Florida: Chapman; Hall/CRC. <https://bookdown.org/yihui/rmarkdown-cookbook>.
- Yackulic, Charles B., Michael Dodrill, Maria Dzul, Jamie S. Sanderlin, and Janice A. Reid. 2020. "A Need for Speed in Bayesian Population Models: A Practical Guide to Marginalizing and Recovering Discrete Latent States." *Ecological Applications* 30 (5): e02112. <https://doi.org/10.1002/eap.2112>.
- Zhang, Wen, Jeffrey Pullin, Lyle Gurrin, and Damjan Vukcevic. 2022. "Investigating the Efficiency of Marginalising over Discrete Parameters in Bayesian Computations." *arXiv:2204.06313 [Stat]*. <https://doi.org/10.48550/ARXIV.2204.06313>.

Jeffrey M. Pullin
University of Melbourne
School of Mathematics and Statistics
Melbourne, Australia
ORCID: [0000-0003-3651-5471](https://orcid.org/0000-0003-3651-5471)

Lyle C. Gurrin
University of Melbourne
School of Population and Global Health
Melbourne, Australia
ORCID: [0000-0001-7052-1969](https://orcid.org/0000-0001-7052-1969)

Damjan Vukcevic
Monash University
Department of Econometrics and Business Statistics
Melbourne, Australia
ORCID: [0000-0001-7780-9586](https://orcid.org/0000-0001-7780-9586)
damjan.vukcevic@monash.edu